




# CUDA/GPU Introduction

Nicolas Blin - AI DevTech <[nblin@nvidia.com](mailto:nblin@nvidia.com)>



# NVIDIA


Application  
Frameworks



Platform




NVIDIA  
AI




NVIDIA Omniverse

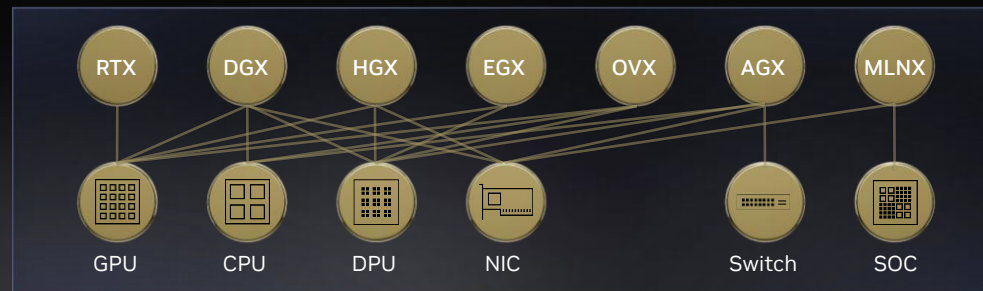
Acceleration  
Libraries



System  
Software



Hardware

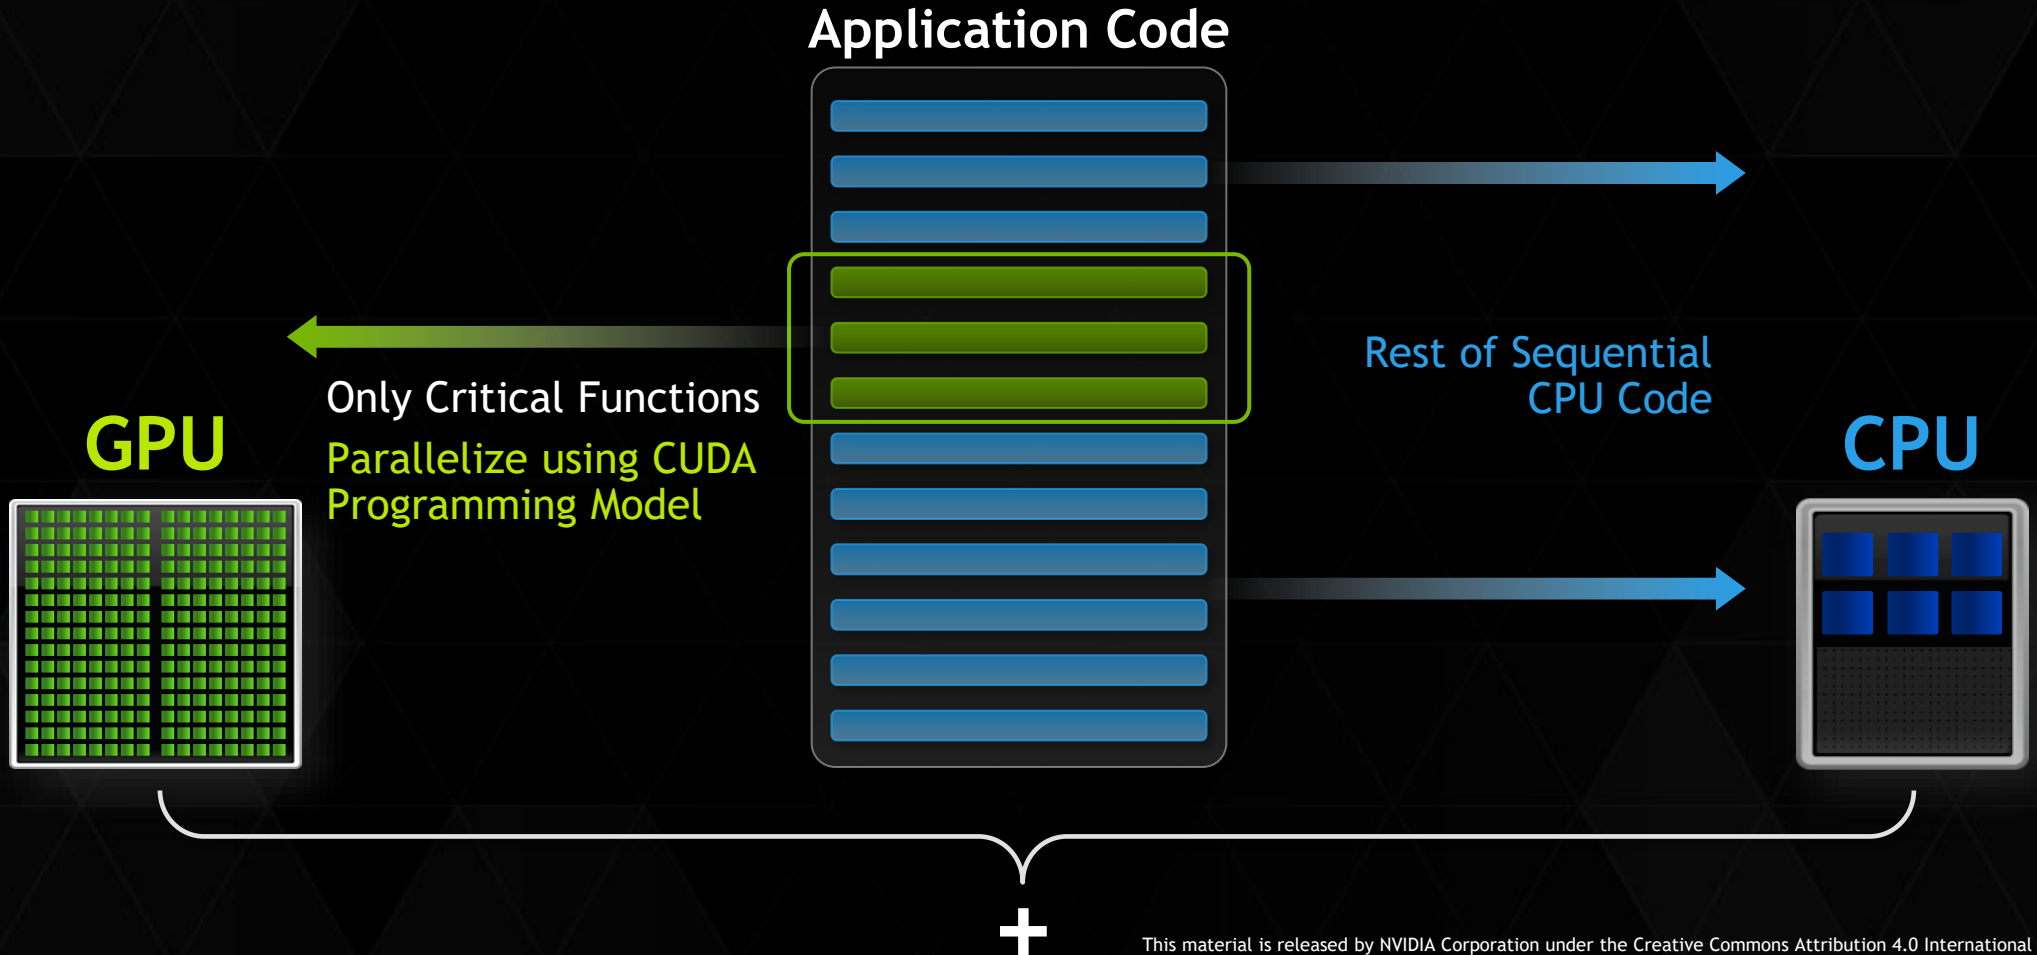


# CUDA C++

## What to expect?

- Basic introduction to GPU Architecture
- GPU Memory and Programming Model
- CUDA C++ programming

# GPU COMPUTING





# CUDA ARCHITECTURE PROGRAMMING MODEL

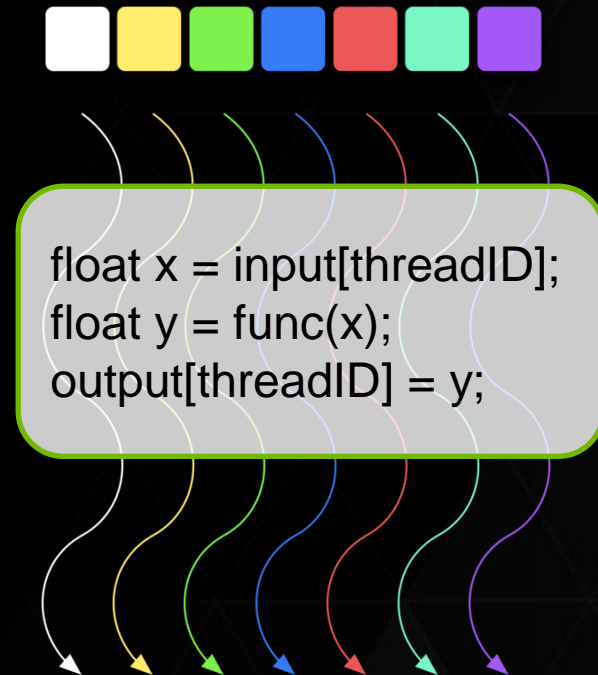
# CUDA KERNELS

- Parallel portion of application: execute as a kernel
  - Entire GPU executes kernel, many threads
- CUDA threads:
  - Lightweight
  - Fast switching
  - Tens of thousands execute simultaneously

CPU	Host	Executes functions
GPU	Device	Executes kernels


# CUDA KERNELS: PARALLEL THREADS

- A **kernel** is a function executed on the GPU
  - Array of threads, in parallel
- All threads execute the same code, can take different paths
  - Each thread has an ID
  - Select input/output data
  - Control decisions






# CUDA KERNELS: SUBDIVIDE INTO BLOCKS






# CUDA KERNELS: SUBDIVIDE INTO BLOCKS




- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**

# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



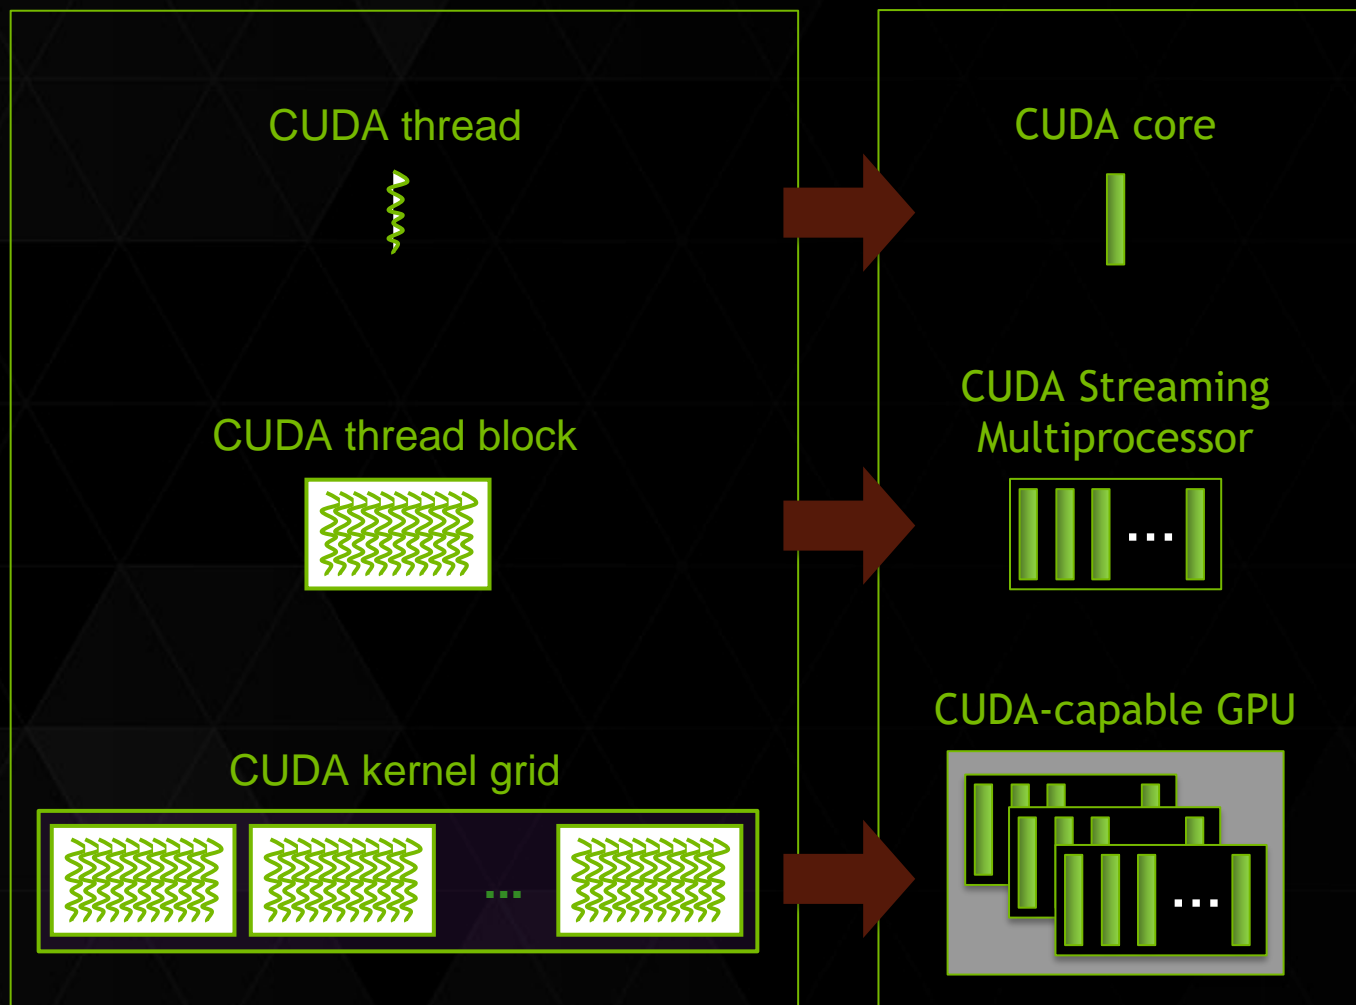
- Threads are grouped into **blocks**
- **Blocks** are grouped into **a grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**

# CUDA KERNELS: SUBDIVIDE INTO BLOCKS



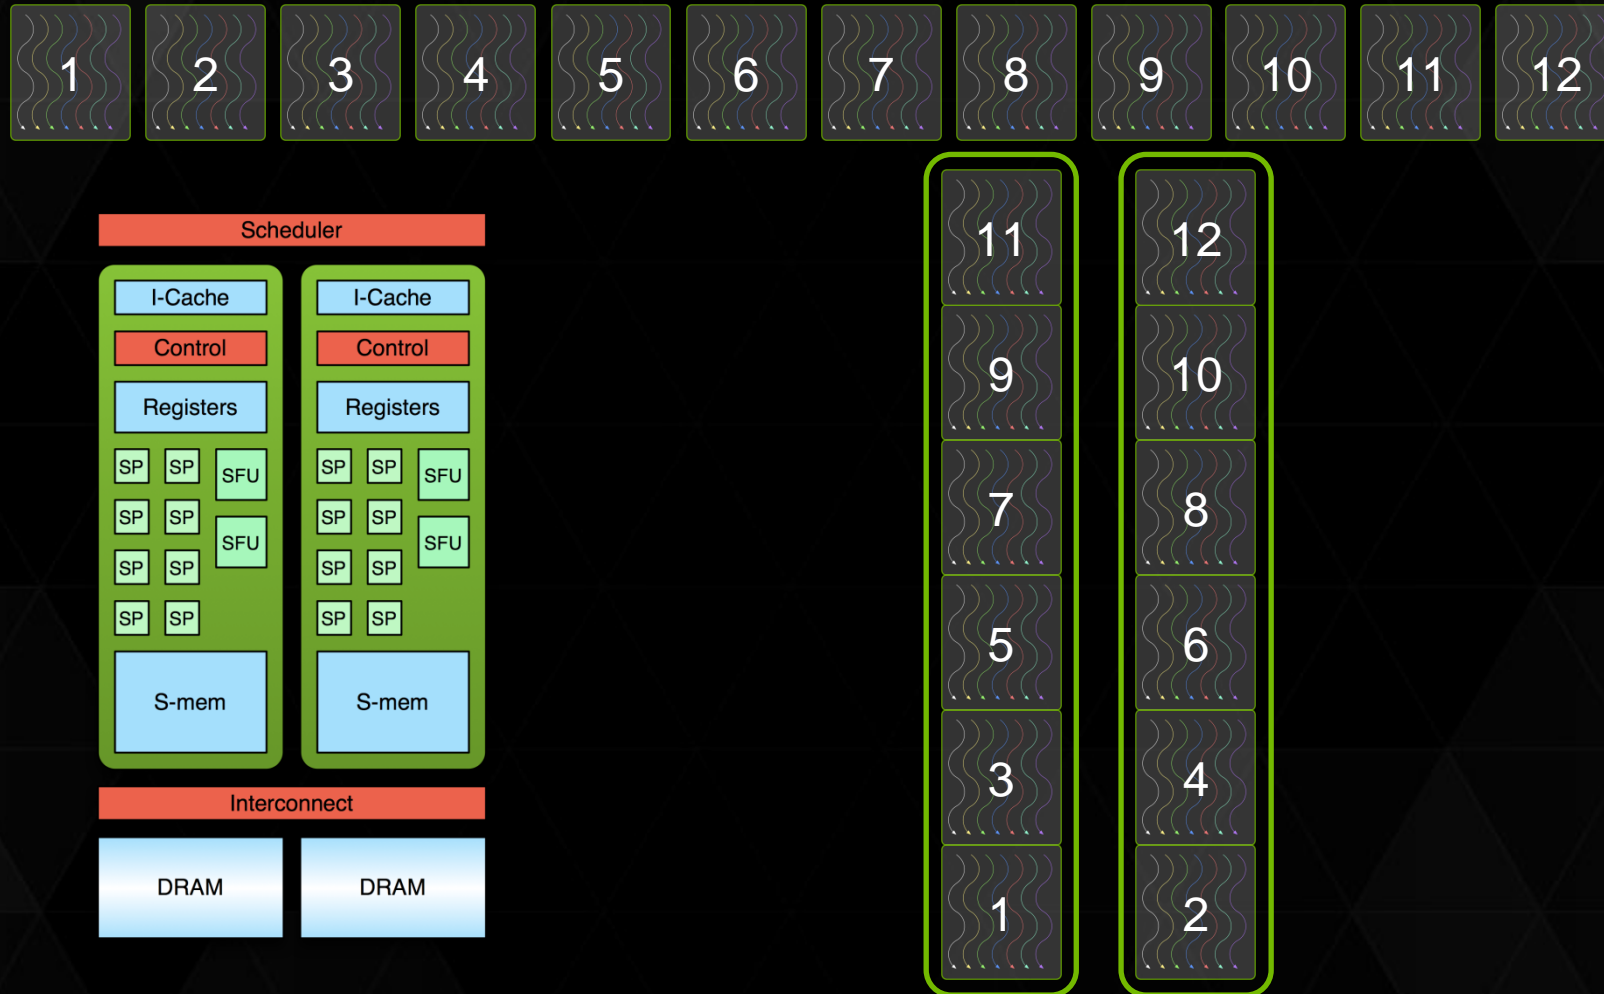
- Threads are grouped into **blocks**
- **Blocks** are grouped into a **grid**
- A **kernel** is executed as a **grid** of **blocks** of **threads**

# KERNEL EXECUTION

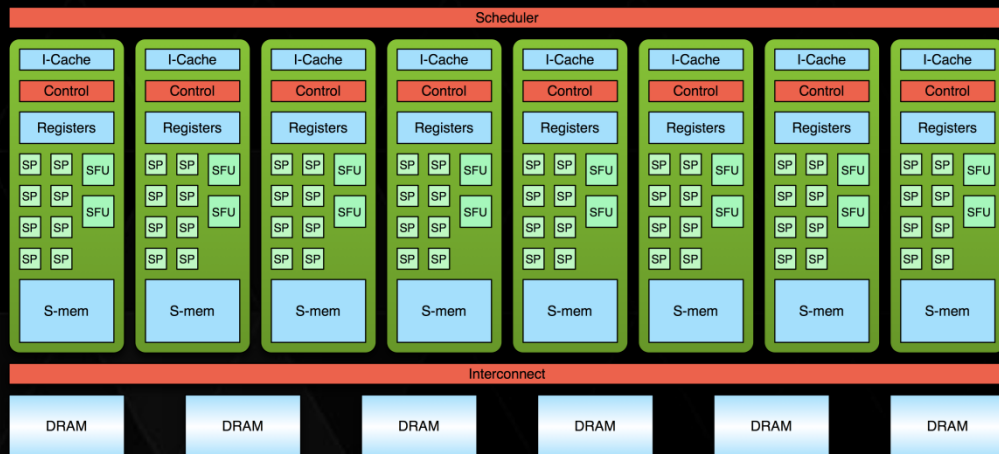


- Each thread is executed by a core
- Each block is executed by one SM and does not migrate
- Several concurrent blocks can reside on one SM depending on the blocks' memory requirements and the SM's memory resources
- Each kernel is executed on one device
- Multiple kernels can execute on a device at one time

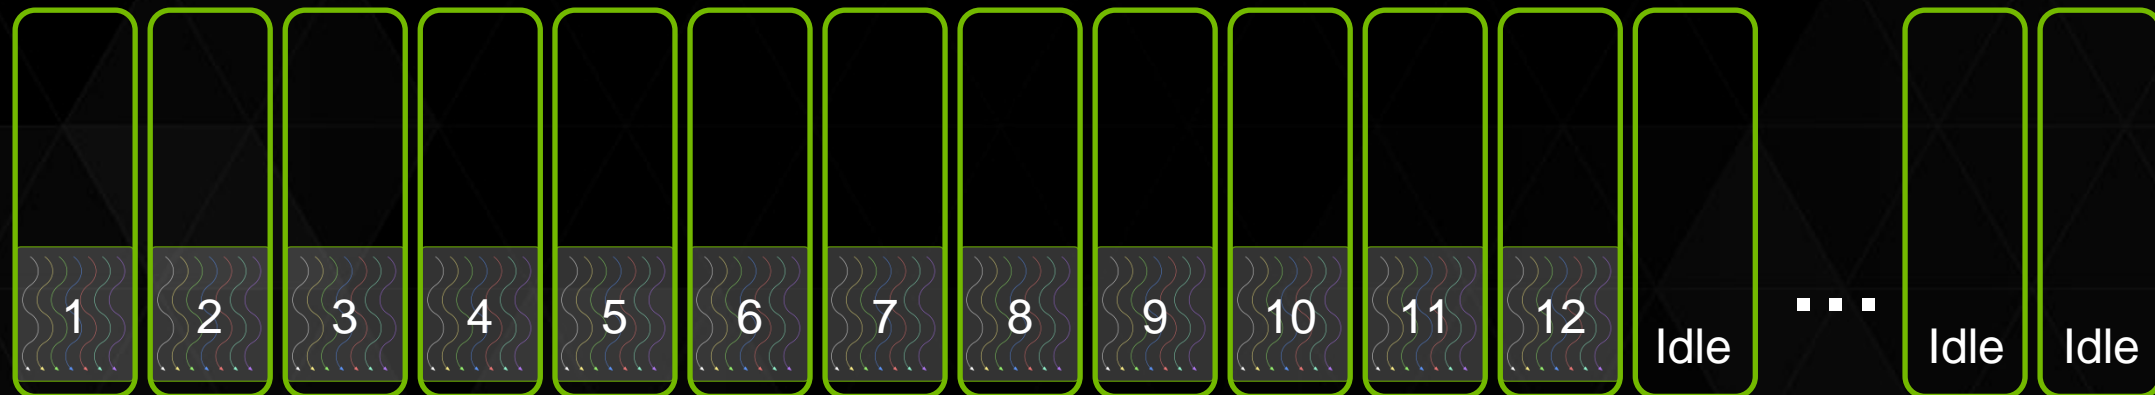
# TRANSPARENT SCALABILITY



# TRANSPARENT SCALABILITY



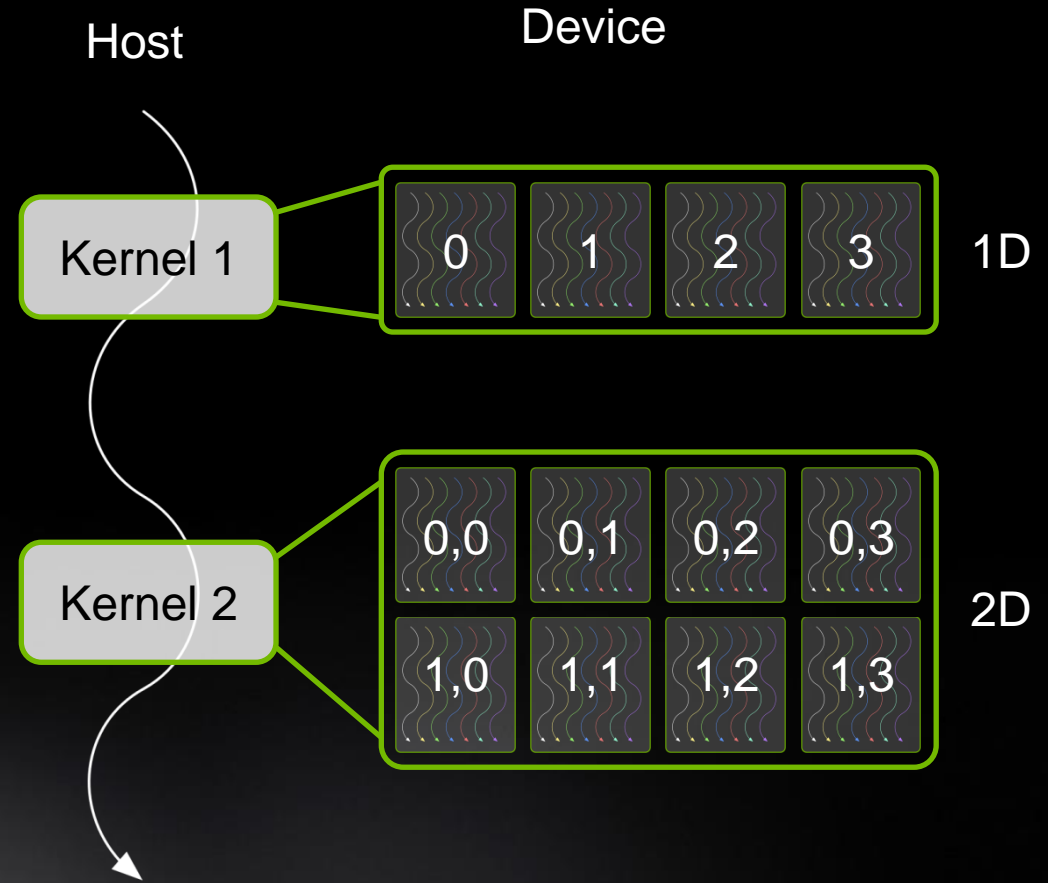
# TRANSPARENT SCALABILITY -





# CUDA PROGRAMMING MODEL - SUMMARY

- A kernel executes as a grid of thread blocks
- A block is a batch of threads
  - Communicate through shared memory
- Each block has a block ID
- Each thread has a thread ID





# CUDA ARCHITECTURE MEMORY MODEL

# GPU ARCHITECTURE

Two Main components

## Global memory

Analogous to RAM in a CPU server

Accessible by both GPU and CPU

48GB with bandwidth currently up to 1 TB/s

## Streaming Multiprocessors (SMs)

SMs perform the actual computations

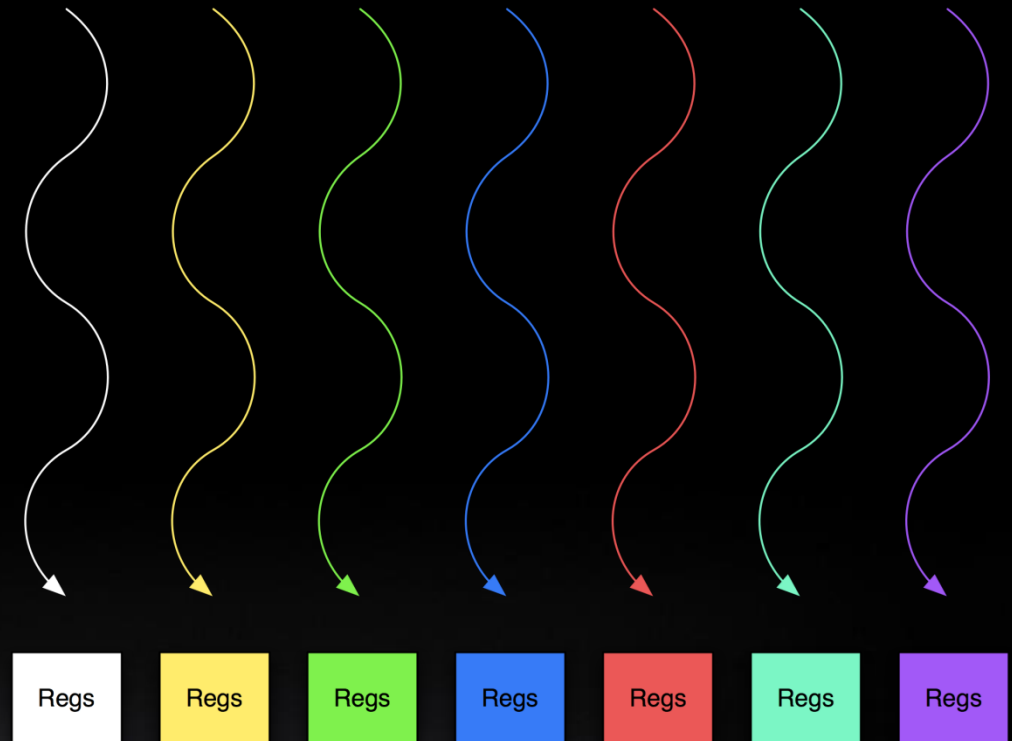
Each SM has its own:

Control units, registers, execution pipelines, caches



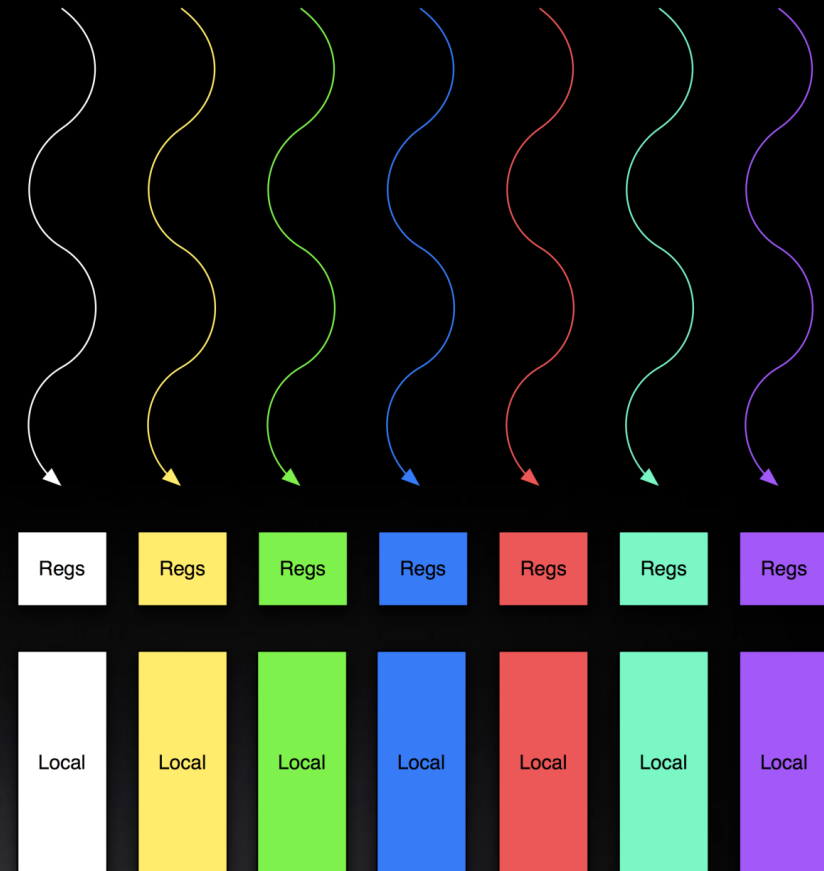
# MEMORY HIERARCHY

- Thread
  - Registers




# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory




# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory
- Block of threads
  - Shared memory



# MEMORY HIERARCHY

- Thread
  - Registers
- Thread
  - Local memory
- Block of threads
  - Shared memory
- All blocks
  - Global memory







# CUDA-C++ INTRODUCTION

# WHAT IS CUDA?

- CUDA Architecture

- Expose general-purpose GPU computing as first-class capability
- Retain traditional DirectX/OpenGL graphics performance


- CUDA C++

- Based on industry-standard C++
- A handful of language extensions to allow heterogeneous programs
- Straightforward APIs to manage devices, memory, etc.

# CUDA C++: THE BASICS

- Terminology
  - *Host* - The CPU and its memory (host memory)
  - *Device* - The GPU and its memory (device memory)

Host



Device



# HELLO, WORLD!

```
int main( void ) {  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

- This basic program is just standard C++ that runs on the *host*
- NVIDIA's compiler (`nvcc`) will not complain about CUDA programs with no *device* code
- At its simplest, CUDA C++ is just C++!

# HELLO, WORLD! WITH DEVICE CODE

```
__global__ void kernel( void ) {  
    printf( "Hello, World!\n" );  
}  
  
int main( void ) {  
    kernel<<<1,1>>>();  
    cudaDeviceSynchronize();  
    return 0;  
}
```

- Two notable additions to the original “Hello, World!”

# HELLO, WORLD! WITH DEVICE CODE

```
__global__ void kernel( void ) {  
}
```

- Keyword `__global__` in CUDA indicates that a function
  - Runs on the device
  - Called from host code
- `nvcc` splits source file into host and device components
  - NVIDIA's compiler handles device functions like `kernel()`
  - Standard host compiler handles host functions like `main()`
    - `g++`
    - Microsoft Visual C++

# NVIDIA HPC SDK

- Comprehensive suite of compilers, libraries, and tools used to GPU accelerate HPC modeling and simulation application
- The NVIDIA HPC SDK includes the new NVIDIA HPC compiler supporting CUDA C and Fortran
  - The command to compile C++ code is 'nvc++'
  - Same compiler will be used for host & device code → more optimization opportunities

```
nvc++ main.cu
```



# HELLO, WORLD! WITH DEVICE CODE

```
int main( void ) {  
    kernel<<< 1, 1 >>>();  
    printf( "Hello, World!\n" );  
    return 0;  
}
```

- Triple angle brackets mark a call from *host* code to *device* code
  - Sometimes called a “kernel launch”
  - We’ll discuss the parameters inside the angle brackets later
- This is all that’s required to execute a function on the GPU!

# A MORE COMPLEX EXAMPLE


- A simple kernel to add two integers:
- As before, `__global__` is a CUDA keyword meaning
  - `add()` will execute on the device ... so `a`, `b`, and `c` must point to device memory
  - How do we allocate memory on the GPU?
  - `add()` will be called from the host

```
__global__ void add( int *a, int *b, int *c ) {  
    *c = *a + *b;  
}
```

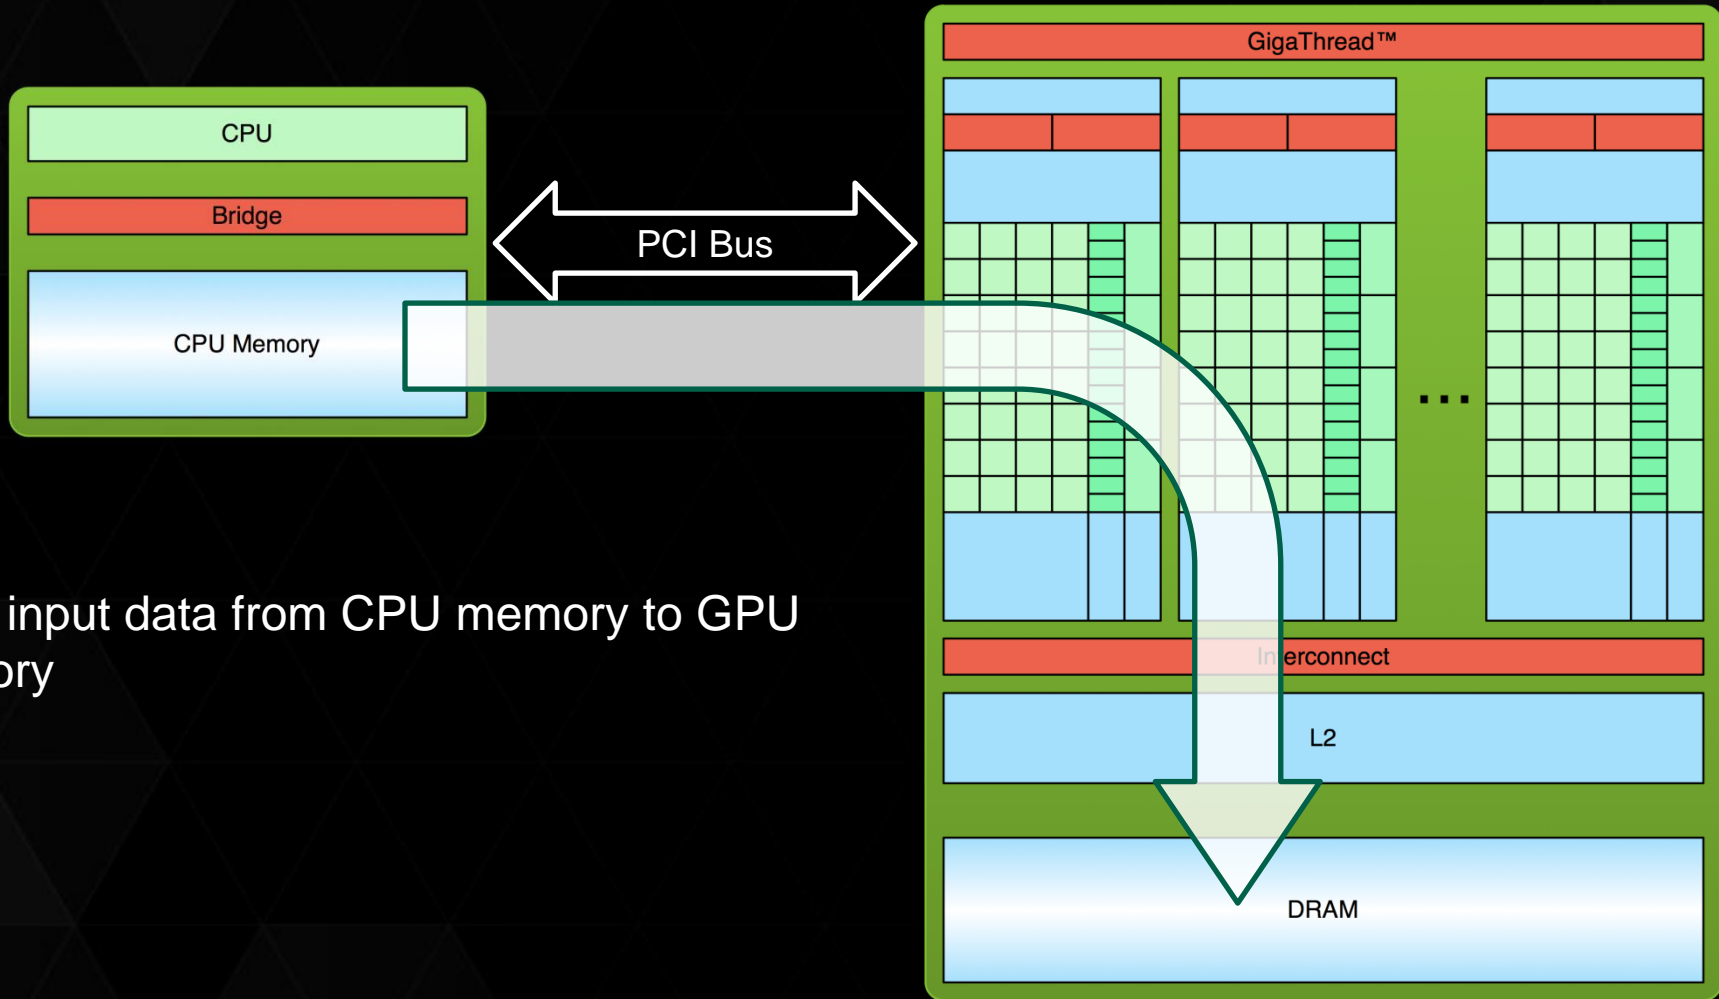
# CPU + GPU

## Physical Diagram

- CPU memory is larger, GPU memory has more bandwidth
- CPU and GPU memory are usually separate, connected by an I/O bus (traditionally PCIe)
- Any data transferred between the CPU and GPU will be handled by the I/O Bus
- The I/O Bus is relatively slow compared to memory bandwidth
- The GPU cannot perform computation until the data is within its memory

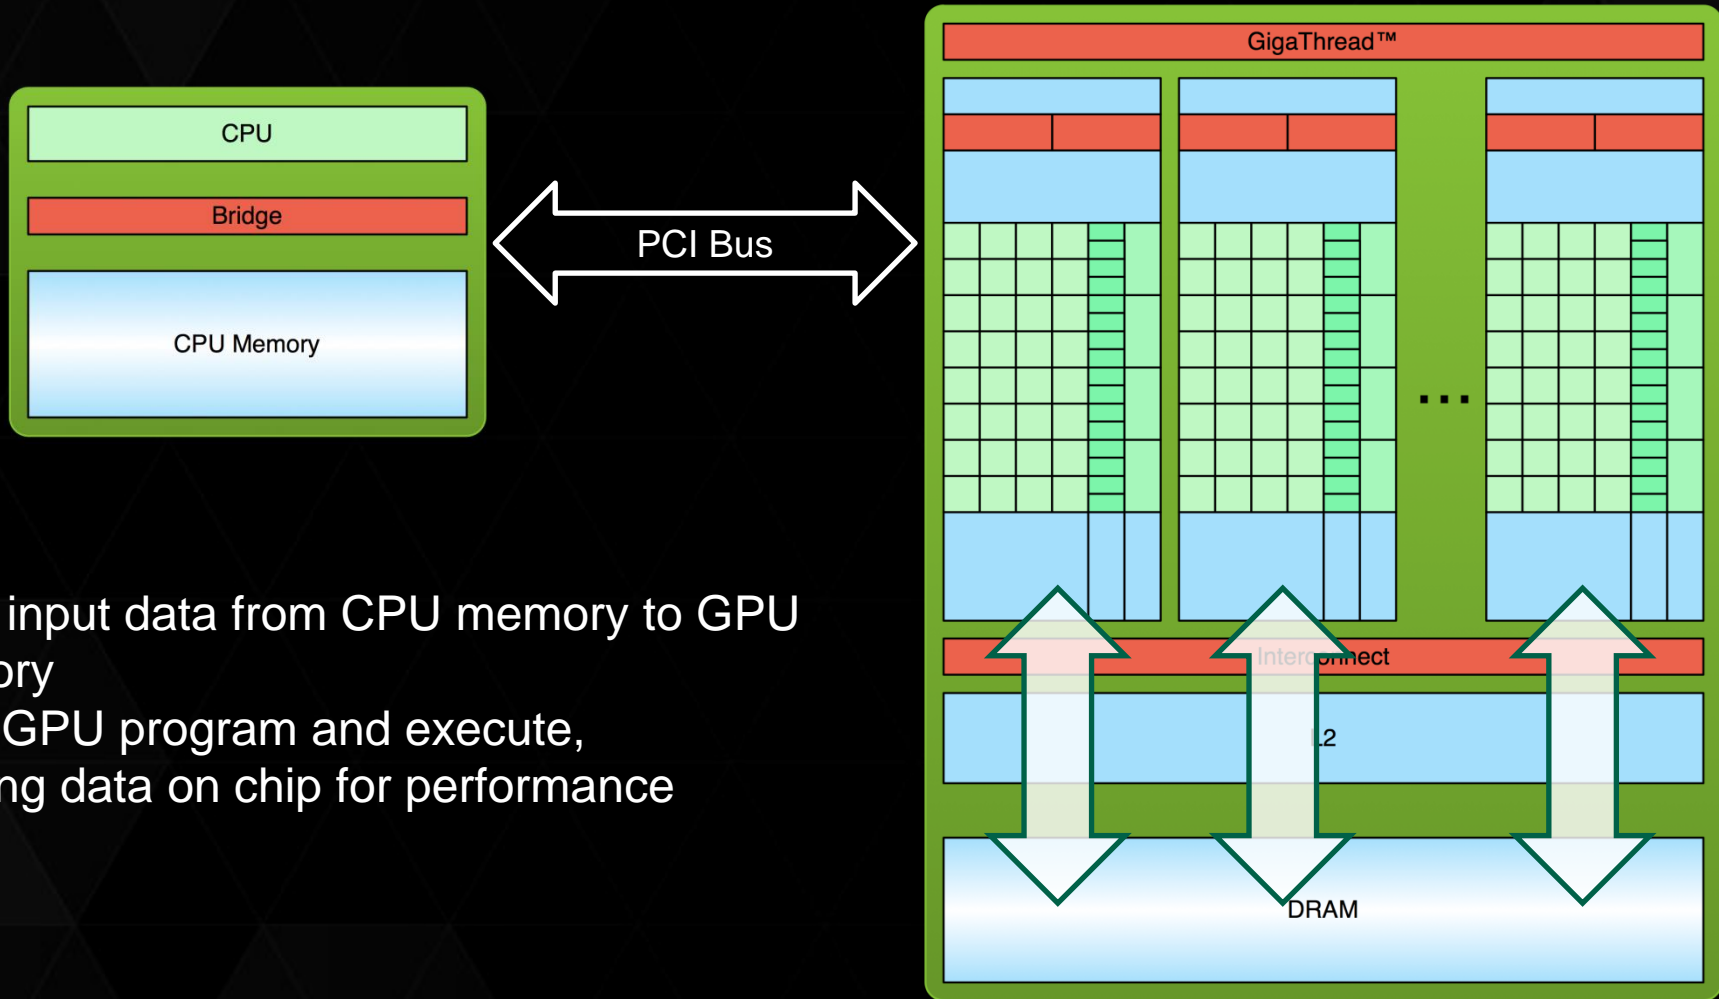


# PROCESSING FLOW - STEP 1

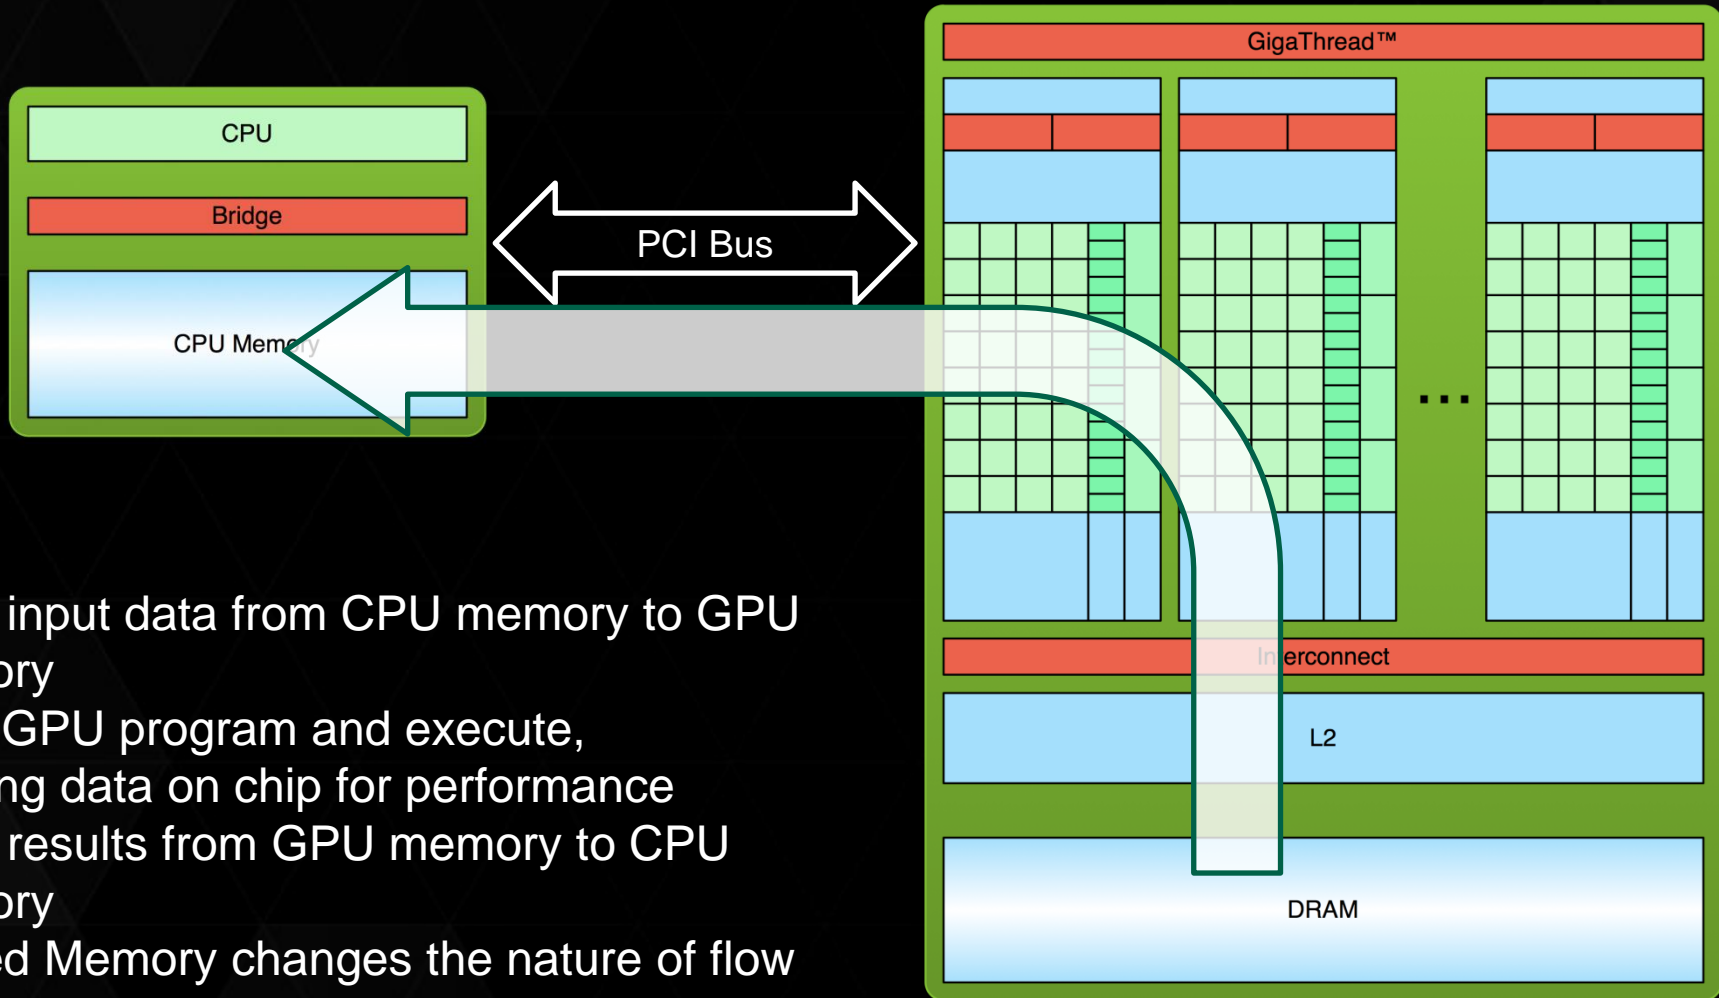


1. Copy input data from CPU memory to GPU memory

# PROCESSING FLOW - STEP 2




# PROCESSING FLOW - STEP 3




1. Copy input data from CPU memory to GPU memory
2. Load GPU program and execute, caching data on chip for performance
3. Copy results from GPU memory to CPU memory
4. Unified Memory changes the nature of flow
  - Some of the basics remains same

# CUDA UNIFIED MEMORY

Simplified Developer Effort



Commonly referred to as  
*"managed memory."*






# MANAGED MEMORY

## Limitations

- Disadvantage:
  - The programmer will almost always be able to get better performance by manually handling data transfers. Just Too Late
  - Memory allocation/deallocation takes longer with managed memory
- Advantage:
  - Handling explicit data transfers between the host and device (CPU and GPU) can be difficult
  - This allows the developer to concentrate on parallelism and think about data movement as an optimization



# CUDA MANAGED MEMORY

## Usefulness

- Use modern C++17 parallel algorithm seamlessly on the GPU:
  - Do all allocations using managed memory
  - Call to all C++ algorithms (*reduce*, *scan*, *transform...*) with *std::par*
  - Compile with `nvc++` and the correct flags

```
nvc++ -stdpar=gpu program.cu -o program
```

# CUDA MANAGED MEMORY

## Usefulness

```
nvc++ -stdpar=gpu program.cu -o program
```

```
int main(int argc, char* argv[]) {  
    std::vector<int> vec(...) // Allocated on the CPU  
    std::reduce(std::par, vec.begin() , vec.end()); // Automatically copied and computed on the GPU  
  
    cpu_computation(vec); // Rest of the complex CPU code kept unchanged  
}
```

# MEMORY MANAGEMENT

- Host and device memory are distinct entities
- Basic CUDA API for dealing with explicit device memory management
  - `cudaMalloc()`, `cudaFree()`, `cudaMemcpy()`
  - Similar to their C equivalents, `malloc()`, `free()`, `memcpy()`
- CUDA API for using Unified memory is
  - C API: `cudaMallocManaged()`, `cudaFree()`
  - \* For this session we will be making use of Unified Memory

# A MORE COMPLEX EXAMPLE: MAIN ( )

```
int main( void ) {  
    int *a, *b, *c;  
    int size = sizeof( int );  
  
    cudaMallocManaged( &a, size );  
    cudaMallocManaged( &b, size );  
    cudaMallocManaged( &c, size );  
  
    add<<< 1, 1 >>>( a, b, c );  
  
    cudaFree( a ); cudaFree( b );  
    cudaFree( c );  
    return 0;  
}
```

# PARALLEL PROGRAMMING IN CUDA

- But wait...GPU computing is about massive parallelism
- So how do we run code in parallel on the device?
- Solution lies in the parameters between the triple angle brackets:

```
add<<< 1, 1 >>>( a, b, c );
```

↓ Number of blocks

```
add<<< N, 1 >>>( a, b, c );
```

- Instead of executing add() once, add() executed N times in parallel

# PARALLEL PROGRAMMING IN CUDA

- With `add()` running in parallel...let's do vector addition
- Terminology: Each parallel invocation of `add()` referred to as a *block*
- Kernel can refer to its block's index with the variable
  - C: `blockIdx.x`
- Each block adds a value from `a[]` and `b[]`, storing the result in `c[]`:
- By using `blockIdx.x` to index arrays, each block handles different indices

# PARALLEL PROGRAMMING IN CUDA

```
__global__ void add( int *a, int *b, int *c) {  
    c[blockIdx.x] = a[blockIdx.x] + b[blockIdx.x];  
}
```

Block 0

```
c[0] = a[0] + b[0];
```

Block 1

```
c[1] = a[1] + b[1];
```

Block 2

```
c[2] = a[2] + b[2];
```

Block 3

```
c[3] = a[3] + b[3];
```

This is what runs in parallel on the device



# THREADS

- Terminology: A block can be split into parallel *threads*
- Let's change vector addition to use parallel threads instead of parallel blocks:
- We use `threadIdx.x` instead of `blockIdx.x` in `add()`

```
__global__ void add( int *a, int *b, int *c) {  
    c[threadIdx.x] = a[threadIdx.x] + b[threadIdx.x];  
}
```

`add<<<N,1>>>>( a, b, c )`




`add<<<1,N>>>>( a, b, c );`

# USING THREADS AND BLOCKS

- We've seen parallel vector addition using
  - Many blocks with 1 thread apiece
  - 1 block with many threads
- Let's adapt vector addition to use lots of *both* blocks and threads
- After using threads and blocks together, we'll talk about *why* threads
- First let's discuss data indexing...

# INDEXING ARRAYS WITH THREADS AND BLOCKS

- No longer as simple as just using `threadIdx.x` or `blockIdx.x` as indices
- To index array with 1 thread per entry (using 8 threads/block)



- If we have  $M$  threads/block, a unique array index for each entry given by

```
int index = threadIdx.x + blockIdx.x * M;
```



Diagram illustrating the mapping of `threadIdx.x` and `blockIdx.x` to the final array index:

```
int index = x + y * width;
```

Arrows indicate the mapping: `threadIdx.x` maps to `x`, `blockIdx.x` maps to `y`, and `M` maps to `width`.

# INDEXING ARRAYS: EXAMPLE

- In this example, the red entry would have an index of 21:



```
int index = threadIdx.x + blockIdx.x * M;  
  
          =      5      +      2      * 8;  
  
          = 21;
```

# ADDITION WITH THREADS AND BLOCKS

- The `blockDim.x` is a built-in variable for threads per block:

```
int index= threadIdx.x + blockIdx.x * blockDim.x
```

- So what changes in `main()` when we use both blocks and threads?

```
__global__ void add( int *a, int *b, int *c ) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
  
    c[index] = a[index] + b[index];  
}
```

# PARALLEL ADDITION (BLOCKS/THREADS): MAIN ( )

```
#define N    (2048*2048)

#define THREADS_PER_BLOCK 512

add<<< N/THREADS_PER_BLOCK, THREADS_PER_BLOCK >>>( a, b, c );
```

# CUDA SPEEDUP : cuOpt

Traveling Salesman Problem : **shortest path to visit a set of points**


Vehicle routing problem (VRP) : multiple vehicles, weight & time constraints...

NP-Hard problem with factorial complexity :

**100 cities → more paths to test than the number of atoms in the universe**



# CUDA SPEEDUP





# REFERENCES

<https://developer.nvidia.com/hpc-sdk>

# Have fun practicing with CUDA

DLI course will allow you to **practice easily and without needing a GPU**:

<https://courses.nvidia.com/courses/course-v1:DLI+T-AC-01+V1/>



# THANK YOU

Any questions?

